AI for Science in Financial Domains
Agent-Governed Discovery Architectures for High-Accountability Professional Work

Algorithmic Trading · Tax Planning · Litigation Strategy · Financial Advice · Investment Banking and M&A

Alejandro Reynoso
Chief Scientist, DEFI Capital

Long Report Version
June 2026

Executive Summary

This long report consolidates five papers that apply the *AI for Science* paradigm to financial and professional advisory domains. The central thesis is that generative artificial intelligence should not be treated as an autonomous decision-maker in high-stakes professional work. Instead, it should be embedded inside a governed discovery architecture in which the model proposes bounded hypotheses, deterministic engines compute outcomes, stress tests expose fragility, professional review agents challenge recommendations, and evaluation agents decide whether the process should continue, stop successfully, or halt after a maximum number of rounds.

The five domains are algorithmic trading, tax planning, civil litigation strategy, household financial advice, and investment banking / M&A. Each domain has different professional risks, different computational engines, and different review standards. Yet all five share the same operating grammar:

$$\text{Generate} \rightarrow \text{Compute} \rightarrow \text{Stress Test} \rightarrow \text{Review} \rightarrow \text{Decide} \rightarrow \text{Iterate.}$$

The contribution of the collection is methodological. The report does not claim that the notebooks constitute production systems, legal advice, tax advice, investment advice, financial advice, or transaction advice. The notebooks use synthetic environments and simplified assumptions. Their purpose is to demonstrate how AI-assisted professional discovery can be made structured, auditable, reproducible, and accountable.

The report also advances a stronger professional standard: in finance-adjacent domains, AI output should be evaluated under both an epistemic burden and an institutional burden. The epistemic burden asks whether a candidate is supported by data, computation, legal authority, evidence, or valuation logic. The institutional burden asks whether the candidate can be explained, supervised, documented, and reviewed by the professionals who remain accountable for the decision. The resulting standard is closer to research governance, fiduciary process, and legal defensibility than to ordinary prompt engineering.

Contents

1. Introduction

1.1. What Is AI4S?

AI for Science, abbreviated AI4S, is the use of artificial intelligence to support disciplined discovery rather than merely automate description, prediction, or communication. In its original scientific setting, AI4S refers to systems that help researchers search enormous spaces of possible explanations, molecules, materials, designs, or experiments [1]. A model may identify a protein structure, propose a molecular candidate, suggest a numerical experiment, infer a physical relationship, or recommend a next measurement. The AlphaFold example is instructive because its significance lies not in fluent explanation, but in the conversion of a difficult biological problem into a structured computational object that could be validated against scientific evidence [2]. Yet the key feature is not that an AI system says something sophisticated. The key feature is that its output enters a loop of testing. A proposal becomes useful only when it can be represented, computed, falsified, compared, and improved.

This distinction matters because AI4S is often misunderstood as a synonym for applying large models to technical problems. The concept is broader and more demanding. AI4S joins generative models, domain data, simulation, mathematical modeling, experimental design, verification, and expert review. It treats AI as a participant in a research workflow, not as a final authority. The model can widen the search space by suggesting candidates that a human team might not have considered. It can compress the time required to organize prior knowledge, translate a vague idea into a formal representation, or propose the next set of tests. But the system remains scientific only if the candidate is exposed to evidence and if failure changes what happens next.

In this report, AI4S is used in that broader methodological sense. The relevant discovery object is not a molecule or a material. It is a professional strategy. In algorithmic trading, the object is a candidate alpha signal or portfolio rule. In tax planning, it is a structure that may reduce tax liability while remaining lawful and audit-defensible. In civil litigation, it is a theory of claims, defenses, settlement posture, and evidentiary development. In household financial advice, it is a plan that organizes goals, risk, liquidity, taxes, and behavior. In investment banking and M&A, it is a potential transaction thesis. These objects are not scientific in the laboratory sense, but they can be handled through a similar discovery grammar: propose, compute, stress, review, decide, and iterate.

The central claim is that high-stakes professional work needs AI4S more than it needs unconstrained automation. A generative model can produce fluent answers to complex questions, but fluency is not the same as validity. A plausible trading story may fail after costs. A confident tax strategy may overlook a legal gate. A compelling litigation narrative may be unsupported by evidence. A financial plan may look rational in a spreadsheet but collapse under household behavior. A deal thesis may look accretive in a base case while failing under integration risk. In each case, the professional danger is not simply error. It is persuasive error: output that sounds competent enough to bypass the discipline that normally protects clients, firms, courts, markets, and boards.

AI4S responds by separating creativity from authorization. The generative component is allowed to propose, recombine, and explore. The deterministic component computes consequences under defined assumptions. The stress-testing component asks what happens when assumptions deteriorate. The professional-review component challenges the output against domain standards. The evaluation component decides whether the process should continue, stop successfully, or halt because the search has reached its limit. This is why the architecture is more important than any single prompt or model. A responsible system is not built around the hope that the model will always be right. It is built around the expectation that many proposals will fail, and that the workflow must make failure visible, cheap, structured, and instructive.

1.2. Areas of Application

The most familiar applications of AI4S appear in the natural and physical sciences. In biology, AI systems can help predict structures, analyze sequences, design proteins, identify biomarkers, and prioritize experiments. In chemistry, they can search chemical space, suggest synthetic routes, estimate properties, and screen molecular candidates. In materials science, they can help identify compounds with desired properties such as conductivity, durability, lightness, or thermal stability. In physics and astronomy, AI can help classify signals, accelerate simulation, detect anomalies, and support inverse problems where researchers infer hidden causes from observed data. In climate and earth systems, AI can help integrate large data streams, emulate expensive simulations, and identify patterns that would be difficult to detect manually.

AI4S also applies to engineering and complex design. Aerospace, robotics, energy systems, manufacturing, and infrastructure all involve large design spaces with many constraints. A design may be aesthetically appealing or technically novel and still fail because of weight, cost, reliability, safety, maintenance, or regulatory limits. AI4S is useful where candidate generation must be paired with simulation and validation. The model may propose a design family, while engineering tools evaluate stress, efficiency, resilience, manufacturability, and lifecycle performance. The discovery loop remains the same: the system learns not only from winners but also from failed candidates that reveal which regions of the design space are infeasible.

Medicine and health care provide another important area of application, though with especially strict limits. AI can help organize medical literature, discover drug candidates, segment images, identify trial cohorts, predict deterioration, or recommend diagnostic workups for review. But medical work illustrates the danger of confusing assistance with authority. A model-generated clinical suggestion must be constrained by evidence, patient context, regulation, professional judgment, and safety review. The AI4S pattern is valuable precisely because it does not treat a model output as care. It treats it as a candidate to be checked through a clinical and institutional process.

The same pattern can extend beyond the laboratory into social science, law, policy, education, and public administration. These domains often lack the clean experimental conditions of chemistry or physics, but they still involve hypotheses, data, constraints, interventions, and evaluation. A policy proposal can be simulated, stress-tested, compared with historical evidence, reviewed for legal authority, and revised when it fails. An educational intervention can be designed, evaluated, and adapted. A legal argument can be checked against facts, precedent, procedure, and ethics. In these areas, AI4S does not mean pretending that professional judgment is a laboratory experiment. It means importing the discipline of hypothesis formation, testing, review, and revision into knowledge work that is otherwise vulnerable to narrative confidence.

Professional services are therefore a natural next frontier. Accounting, legal advice, wealth management, investment research, corporate finance, insurance, risk management, and compliance all require reasoning under uncertainty. They also generate records, models, classifications, rules, forecasts, scenarios, and review trails. These characteristics make them suitable for AI4S architectures. The model can help produce candidate analyses, but every candidate must pass through calculation, suitability, legal, ethical, and institutional gates. The result is not a replacement for the professional. It is an augmented research environment in which the professional can see more alternatives, test more assumptions, and document more of the reasoning process.

1.3. Why Finance?

Finance is an especially important domain for AI4S because it combines data abundance with high accountability. Financial decisions are often measurable, modelable, and repeatable, but they are also consequential. A trading system can lose capital quickly. A tax structure can expose a client to penalties or audit risk. A financial plan can shape a household's retirement, insurance, education, and liquidity choices. A legal strategy can affect settlement value and legal rights. An acquisition can reshape a company, alter employment, change competitive dynamics, and consume large amounts of shareholder capital. In finance, the gap between plausible analysis and responsible advice can be large.

The field also has a long tradition of formal modeling. Portfolio theory, asset pricing, risk measurement, valuation, accounting, credit analysis, and deal modeling all rely on mathematical and computational tools [3, 4]. This makes finance attractive for AI4S because many candidate ideas can be translated into structured inputs and tested. A trading hypothesis can become a signal definition and backtest. A tax idea can become a calculation under specified assumptions. A financial plan can become a cash-flow projection and Monte Carlo simulation. A deal thesis can become a valuation model, financing plan, accretion-dilution analysis, and synergy case [5, 6]. The presence of deterministic engines gives the AI4S loop something concrete to do after generation.

At the same time, finance is not merely quantitative. It is institutional, legal, behavioral, strategic, and ethical. The quality of a recommendation depends not only on expected return, tax savings, valuation, or modeled probability. It also depends on suitability, fiduciary duty, disclosure, conflicts, liquidity, market impact, client understanding, board process, and regulatory context. A purely computational answer may be incomplete, while a purely narrative answer may be unsafe. AI4S is useful because it can combine structured computation with professional review. The architecture forces the system to ask not only whether an idea looks attractive, but also whether it is permissible, robust, explainable, documented, and appropriate for the specific mandate.

Finance is also adversarial and adaptive. Markets change when participants exploit patterns. Tax authorities challenge aggressive structures. Opposing counsel attacks weak litigation theories. Clients behave differently under stress than they do in intake forms. Sellers, buyers, lenders, regulators, and boards react strategically in M&A. This means that a candidate recommendation cannot be judged only in a base case. It must be exposed to stress. What happens if transaction costs rise, volatility changes, liquidity disappears, audit scrutiny increases, settlement leverage weakens, the client loses income, or synergies arrive late? Stress testing is therefore not an optional add-on. It is the mechanism that turns a clever idea into a risk-aware professional object.

The five domains in this report were selected because they represent different versions of this same problem. Algorithmic trading tests whether language-model creativity can be converted into bounded alpha hypotheses and rejected when the evidence fails. Tax planning tests whether a generated strategy can be filtered by legal rules, suitability, audit posture, and CPA-style review. Litigation strategy tests whether an AI-generated theory can respect evidence, ethics, procedure, and settlement reality. Financial advice tests whether recommendations can be aligned with household goals, risk tolerance, liquidity needs, taxes, and behavioral constraints. Investment banking and M&A tests whether a transaction idea can survive valuation, financing, synergy, diligence, regulatory, and board-readiness review.

The purpose is not to claim that these simplified notebooks are production systems. They are demonstrations of method. They show how an AI-assisted workflow can be structured so that the model's creative contribution is captured without granting the model unchecked professional authority. They also show how failures can become useful. A rejected trading strategy, fragile tax plan, weak litigation theory, unsuitable financial recommendation, or non-board-ready deal is not merely a negative result. It is evidence that can guide the next search. In this sense, finance becomes a laboratory for accountable AI4S: a setting where ideas are plentiful, risks are explicit, validation is necessary, and governance cannot be separated from usefulness.

1.4. Argument of This Report

The report argues that AI4S should be understood as a governance architecture for professional discovery. Its value does not come from replacing judgment with automation. Its value comes from making judgment more testable, transparent, and reviewable. The model expands the set of possible strategies. The computational engine translates proposals into consequences. Stress tests reveal fragility. Review agents represent professional skepticism. Evaluation agents decide whether the process has produced a candidate worth preserving or a failure worth learning from. The architecture is iterative because professional discovery is iterative.

This argument matters because many discussions of AI in finance focus on two extremes. One extreme imagines autonomous expert systems that produce finished decisions. The other treats generative AI as a drafting tool that merely writes memos faster. The AI4S view is different. It asks how AI can help teams search, compute, challenge, and improve strategies while preserving accountability. It does not deny that language models can be powerful. It insists that their power should be organized inside a process that understands the difference between a proposal and a decision.

The chapters that follow develop this pattern across five domains. Each chapter identifies the domain problem, explains the architecture, describes the computation and review logic, and discusses improvements needed for real-world adaptation. The cross-domain synthesis then shows that the same operating grammar appears repeatedly even when the professional risk object changes. The conclusion returns to the central principle: AI may expand the search space, but it may not collapse the review process. In finance and adjacent professional domains, that principle is not a constraint on innovation. It is the condition that makes responsible innovation possible.

The standard of analysis is therefore deliberately dual. From finance, the report imports attention to uncertainty, incentives, market frictions, valuation discipline, risk transfer, and capital allocation. From law, it imports attention to authority, procedure, fiduciary responsibility, evidentiary sufficiency, disclosure, conflicts, and reviewable process. The resulting framework treats professional AI systems as institutional decision systems rather than isolated prediction engines.

The report makes four specific contributions. First, it defines AI4S as an auditable discovery loop rather than a generic use of AI in technical work. Second, it translates that loop into five finance-adjacent domains with distinct risk objects and professional gates. Third, it introduces evaluation burdens that force each generated candidate to become a structured, computable, stress-tested, and reviewable object. Fourth, it shows how failure-context memory can convert rejected outputs into disciplined search information for later rounds.

2. The AI4S Governance Pattern

2.1. The Common Problem

Professional advisory work is rarely a matter of producing a single answer. A trading signal must survive backtesting, costs, liquidity, and regime stress. A tax strategy must be legal, suitable, and audit-defensible. A litigation theory must be ethical, evidence-based, and settlement-aware. A financial plan must respect the client's hierarchy of goals, risk tolerance, liquidity needs, and behavioral constraints. An M&A target must be strategically compelling, financeable, synergy-defensible, accretive, and board-ready.

In all of these cases, the risk of generative AI is not merely that it may be wrong. The deeper risk is that it may be persuasive. It may produce fluent, confident, and plausible narratives that create the appearance of professional reasoning without the underlying calculation, verification, or accountability that professional work requires.

The problem is therefore one of institutional epistemology: how can a firm know that a machine-assisted recommendation is warranted? In finance, the answer cannot be reduced to predictive accuracy because many errors arise from model risk, incentives, timing, liquidity, and omitted constraints. In law, the answer cannot be reduced to rhetorical plausibility because professional judgment depends on authority, procedure, facts, burdens, ethics, and client authorization. The architecture must therefore generate evidence about the process, not merely content about the answer.

2.2. The AI4S Translation

The AI4S paradigm treats intelligence as part of a discovery loop. In scientific settings, this means proposing hypotheses, testing them against experiments, learning from failures, and updating the next search direction. In financial and professional domains, the same logic can be adapted:

AI4S Scientific Concept	**Professional Advisory Equivalent**
Observation	Case facts, client profile, market data, mandate, or tax profile
Hypothesis	Strategy, legal theory, alpha signal, wealth plan, or acquisition target
Experiment	Backtest, tax simulation, precedent simulation, Monte Carlo run, or deal model
Validation	Legal gate, suitability gate, ethics gate, regulatory gate, or stress test
Peer review	CPA, attorney, CFP, MD, or research-governance review
Iteration	Failure-context injection into the next cycle

2.3. The Shared Architecture

Across the five papers, the general architecture is:

Structured Domain Universe
$\downarrow$
Case / Client / Market / Mandate Profile
$\downarrow$
LLM Candidate Generation
$\downarrow$
Parser and Validation Layer
$\downarrow$
Domain-Specific Deterministic Engine
$\downarrow$
Stress Testing
$\downarrow$
Professional Review Agent
$\downarrow$
Evaluation Agent
$\downarrow$
Failure Context Injection.

This architecture separates creativity from authorization. The LLM proposes. Code computes. Stress tests challenge. Professional review scrutinizes. Governance decides.

2.4. Evaluation Burdens

A candidate generated by an AI4S workflow should satisfy four cumulative burdens before it can be treated as professionally meaningful. First, it must satisfy a *representational burden*: the output must be reducible to a structured object with defined inputs, assumptions, scope, and exclusion criteria. Second, it must satisfy a *computational burden*: the relevant consequences must be calculated by a deterministic or otherwise validated engine rather than inferred from prose. Third, it must satisfy a *stress burden*: the candidate must be exposed to adverse facts, parameter changes, market conditions, procedural constraints, or behavioral shocks. Fourth, it must satisfy a *review burden*: a professional standard must determine whether the candidate is suitable, lawful, ethical, financeable, admissible, board-ready, or otherwise fit for the mandate.

The governing object is therefore not the model output alone. It is the full evidentiary record generated by the workflow. In formal terms, a candidate h is not accepted because the model assigns it high narrative plausibility. It is accepted only if the record

$$\mathcal{R}(h) = \{\text{inputs}, \text{assumptions}, \text{computation}, \text{stress tests}, \text{review}, \text{decision}\}$$

meets the applicable domain threshold. This formulation is intentionally conservative. It treats AI as a source of hypotheses and organizing intelligence, while reserving professional authority for procedures that can be inspected, challenged, and repeated.

3. AI4S Applied to Algorithmic Trading

3.1. Domain Problem

Algorithmic trading is not simply the search for a clever signal. It is a disciplined discovery problem under noise, uncertainty, regime instability, transaction costs, liquidity limits, crowding effects, and overfitting risk. A superficially attractive signal can appear persuasive in natural language, perform well in a narrow historical sample, and still fail when exposed to realistic trading constraints. In this domain, the danger is not only that the model may be wrong. The deeper danger is that it may produce a confident story of alpha without the evidence required to support it.

At a doctoral research standard, the relevant question is not whether an AI system can narrate a market anomaly. It is whether the anomaly survives a research design that controls for look-ahead bias, survivorship bias, multiple testing, transaction costs, factor exposure, crowding, capacity, and post-discovery decay. A language model can help generate candidate mechanisms, but the burden of proof remains empirical and institutional. The strategy must be precise enough to be replicated by another researcher, rejected by an adverse reviewer, and withheld from deployment unless the evidence justifies the risk.

The problem is therefore to convert AI-generated trading ideas into testable research hypotheses. A trading hypothesis must specify a universe, feature definition, signal direction, portfolio construction rule, holding period, risk control, and rejection condition. If the proposal cannot be parsed, computed, or tested, it is not yet a research object. The objective is not autonomous trading. The objective is controlled alpha discovery, where language-model creativity expands the search space while deterministic tests determine whether a candidate deserves further attention.

This framing is especially important because markets are adaptive. A pattern that seems meaningful may be a statistical artifact, a temporary regime effect, a liquidity premium, or a hidden exposure to common risk factors. The architecture must therefore make failure inexpensive, visible, and informative. A rejected strategy is not wasted output; it is evidence about where not to search next.

That requirement changes the meaning of model success. The model is not rewarded for sounding original, confident, or intuitively appealing. It is rewarded only when its proposal can be translated into a constrained object that the research system can evaluate. A useful hypothesis must identify what is being measured, why the feature might predict returns, when the signal is formed, when the position is entered, how exposure is sized, and what empirical result would falsify the idea. This makes the process closer to experimental design than discretionary commentary. It also protects the research team from the common failure mode in which a broad market story is mistaken for a tradable signal.

The discovery problem is also cumulative. Because many candidate signals will fail, the system must preserve the reason for failure rather than merely discard the output. A rejected hypothesis may reveal that a feature is too noisy, that the holding period is too short, that turnover overwhelms gross alpha, or that the apparent edge depends on an unstable regime. Those diagnoses become inputs to the next round of search. The goal is therefore not to eliminate failure, but to make failure cheap, structured, and useful.

3.2. Governed Architecture

The algorithmic-trading architecture follows the AI4S loop of proposing, computing, stressing, reviewing, deciding, and iterating. The LLM generates bounded alpha hypotheses, but it does not trade, validate itself, or decide that an idea is profitable. Its role is hypothesis generation. A parser then converts the proposal into a structured representation that can be evaluated by code. Ideas that fall outside the permitted vocabulary, reference unavailable data, or fail validation are rejected before backtesting.

The deterministic engine builds features, constructs tradable portfolios, applies transaction costs, and computes performance metrics. Let $r_{i,t}$ denote the return of asset i at time t, and let $s_{i,t}$ denote the signal score generated by a candidate hypothesis. The system forms portfolio weights $w_{i,t}$ as a bounded function of the signal:

$$w_{i,t} = f(s_{i,t}),$$

subject to leverage, concentration, liquidity, turnover, and feasibility constraints. The portfolio return is:

$$R_{p,t} = \sum_i w_{i,t-1} r_{i,t} - TC_t,$$

where TC_t represents transaction costs. A candidate is evaluated through Sharpe ratio, drawdown, turnover, hit rate, stability, capacity, and stress robustness.

After the baseline backtest, the system challenges the strategy under adverse scenarios such as volatility spikes, liquidity deterioration, crowding decay, cost increases, and market-regime shifts. The fully auditable version extends this architecture by preserving every material research artifact: prompt, hypothesis, parser decision, backtest result, stress result, paper-trading result, reviewer comment, and evaluation decision. This turns the notebook from a demonstration of model output into a reproducible research record.

An evaluation agent governs continuation. It stores failed hypotheses, records reasons for rejection, and decides whether the system should continue searching, stop successfully, or halt after the maximum number of rounds. The governing principle is simple: the model may suggest an alpha hypothesis, but only the controlled research process may decide whether that hypothesis survives.

The parser and validation layer are critical because they enforce the boundary between natural-language creativity and executable research. A candidate that references unavailable fields, ambiguous timing, impossible orders, unsupported instruments, or undefined risk controls fails before consuming backtest resources. This design prevents the LLM from smuggling vague discretionary judgment into the deterministic engine. It also ensures that every surviving strategy can be compared against the same research protocol.

The research record should also preserve negative controls. A strong trading workflow should compare the proposed signal against naive benchmarks, permuted labels, delayed signals, sector-neutral variants, and known factor exposures. If an apparently successful signal disappears under these controls, the correct conclusion is not that the model failed creatively; it is that the research process succeeded by rejecting a false discovery.

The iteration layer then turns the research process into a governed search. The evaluation agent does not merely rank Sharpe ratios. It reads performance, stress robustness, turnover, capacity, implementation burden, and reviewer objections together. If the best candidate fails because costs dominate returns, the next prompt can emphasize lower-turnover hypotheses. If a strategy fails under regime stress, the next cycle can ask for signals with explicit regime conditioning. In this way, the architecture operationalizes the AI4S pattern: generate a hypothesis, compute evidence, expose fragility, record the lesson, and search again under tighter constraints.

3.3. Demonstrated Contribution

The main result is not a claim that any specific trading strategy is profitable. The main result is methodological. The notebook shows how generative AI can be used as a disciplined research partner rather than a trading oracle. The LLM expands the space of candidate ideas, while deterministic code preserves analytical integrity.

The architecture creates a workflow in which narrative ideas become computable objects. Weak ideas fail in parsing or backtesting. Fragile ideas fail under stress. Ideas that require unavailable data, unrealistic execution, excessive turnover, or unbounded leverage do not advance. This prevents fluency from being mistaken for evidence.

The system also demonstrates the value of structured failure. Each rejected hypothesis contributes to the next round by clarifying which signal families, assumptions, or constraints proved inadequate. In that sense, the research loop does not merely search for winners. It accumulates a disciplined map of failures, making the discovery process more auditable and less dependent on intuition.

A second result is that the workflow clarifies what counts as evidence. In ordinary AI use, a fluent explanation can create the impression that the strategy has already been reasoned through. In this architecture, explanation is only the beginning of the process. The evidence comes from reproducible feature construction, explicit portfolio rules, transaction-cost accounting, baseline performance, adverse-scenario performance, and review notes. The system therefore gives the researcher a complete trail from idea to rejection or provisional acceptance.

The prototype also shows why governance must be embedded inside the research loop rather than appended at the end. If the model were allowed to generate an idea and then write its own justification, the system would merely automate persuasion. By contrast, the notebook separates roles so that each stage constrains the next one. The LLM supplies diversity of hypotheses, the parser supplies discipline, the backtester supplies measurement, the stress tests supply skepticism, and the evaluation agent supplies memory. The resulting process is slower than simply asking for trade ideas, but it is more suitable for high-accountability research because it makes uncertainty visible.

This matters because an institutional research process must be able to explain why a candidate advanced, why it failed, or why it was stopped before deployment. The notebook therefore produces governance value even when no trade is approved. It demonstrates a repeatable pattern for controlling AI-assisted ideation, measuring ideas under common rules, and preserving enough context for later review by researchers, risk managers, and compliance teams.

3.4. Limitations and Research Extensions

The most important improvement would be to replace the synthetic market environment with institutional-quality historical data. A stronger platform would use survivorship-bias-free universes, point-in-time fundamentals, realistic execution assumptions, borrow costs, slippage models, liquidity constraints, corporate-action handling, and capacity analysis. The backtesting engine could also be expanded to include walk-forward validation, regime classification, factor exposure decomposition, feature-decay analysis, and portfolio-construction constraints.

A second improvement would be to connect the notebook to live paper-trading monitoring. This would allow the system to compare research expectations with out-of-sample behavior before any production use. Paper trading would be especially useful for measuring signal decay, turnover, realized costs, capacity limits, and operational fragility.

Finally, the governance layer could be strengthened into a formal model-risk and research-committee process. Every prompt, model version, output, code version, dataset version, and evaluation decision should be logged. A production-grade version would require human portfolio-manager review, compliance review, risk approval, deployment controls, and clear boundaries preventing the LLM from becoming an execution authority.

The platform could also benefit from stronger statistical controls. Candidate strategies should be evaluated against multiple-testing risk, data-mining bias, benchmark factor exposures, sector and country tilts, and sensitivity to rebalance calendars. A realistic implementation would compare each candidate with simple alternatives, require out-of-sample confirmation, and penalize fragile parameter choices. It should also distinguish between economic intuition and statistical artifact by requiring the research record to explain why a signal might persist after costs and competition.

Operationally, the system should include a promotion pathway from hypothesis to sandbox research, then to paper trading, then to limited capital, and finally to production only after explicit approval. Each stage should have separate thresholds for evidence, risk, monitoring, and shutdown. The LLM should remain outside trade execution and should not be able to alter live code, broker settings, or risk limits. This would preserve the core governance principle while allowing the discovery loop to become useful inside an institutional research process.

The human-review process could also be made more explicit. Research leads should be able to override the evaluation agent, but only with a recorded rationale. Conversely, the system should be able to force escalation when a strategy depends on unusual data, high turnover, concentrated exposure, or unstable market regimes. These controls would make the architecture more credible as a supervised research pipeline rather than a notebook experiment.

3.5. Reference Materials

The companion paper and notebooks for this domain are:

- **Paper:** AI4S ALGO TRADING.
- **Standard notebook:** COLAB ALGO TRADING.
- **Fully auditable notebook:** COLAB ALGO TRADING FULLY AUDITABLE.

4. AI4S Applied to Tax Planning and Audit Governance

4.1. Domain Problem

Tax planning is a high-stakes professional domain where a strategy can be numerically attractive but legally weak, audit-fragile, unsuitable for the taxpayer, or inconsistent with professional standards. A simulated reduction in tax liability is not enough. A strategy must be valid, explainable, documentable, suitable, and defensible under audit. The risk of generative AI is that it may suggest tax-saving ideas that sound plausible but are not grounded in a valid computational or legal framework.

At a higher professional standard, the relevant issue is not tax minimization in isolation. It is whether a position could be taken, documented, explained, and defended under the applicable legal authorities and professional duties. The system must distinguish a computationally favorable scenario from a supportable tax position. That distinction is central because tax advice depends on facts, authority, timing, taxpayer intent, recordkeeping, disclosure, and the practitioner's willingness to sign or defend the work.

The domain problem is therefore to convert tax strategy generation into a governed discovery process. The LLM may propose planning ideas, but those ideas must be separated from tax computation and professional judgment. The system must ask: Is the strategy legally available? Does it apply to this taxpayer? Does it survive stress conditions? Is the documentation burden reasonable? Would a CPA be willing to defend it?

This distinction matters because tax planning is not simply optimization. A strategy that lowers tax in one year may increase future exposure, create liquidity problems, trigger audit risk, conflict with client preferences, or depend on facts that cannot be substantiated. The architecture must therefore treat tax savings as one output of a broader review process, not as the sole objective.

This is a particularly important distinction because tax advice is constrained by facts, law, documentation, timing, and professional standards. A planning idea may be unavailable because the taxpayer lacks the required entity structure, because the income category is wrong, because a phase-out removes the benefit, or because the taxpayer cannot maintain the records needed to support the position. Even when a strategy is technically available, it may be unsuitable if it creates cash-flow pressure, increases future taxable income, or relies on aggressive assumptions that the client would not defend.

The AI4S framing therefore treats each planning idea as a hypothesis about the taxpayer's facts and the applicable rules. The hypothesis must be converted into structured inputs, computed by a deterministic engine, challenged by legal and professional gates, and preserved in an audit trail. The system is not trying to maximize a tax-savings number in isolation. It is trying to identify strategies that are numerically meaningful, legally available, client-specific, documentable, and reviewable. That is why the architecture separates suggestion from computation and computation from authorization.

4.2. Governed Architecture

The tax-planning architecture begins with a synthetic tax code and a taxpayer profile. The deterministic tax engine computes the baseline liability before any AI-generated strategy is considered. The LLM then proposes candidate strategies using a bounded vocabulary. These candidates pass through a legal validity gate that excludes ideas that are unavailable, unsupported, inapplicable, or inconsistent with the taxpayer profile.

The legal gate should be understood as a rule-of-law constraint, not a stylistic review. It asks whether the strategy is authorized by the domain universe, whether the taxpayer facts satisfy the rule predicates, whether the computational result follows from those predicates, and whether the documentation burden is feasible. A strategy that saves money but cannot be substantiated is treated as a professional failure rather than an optimization success.

Only surviving strategies are simulated. Let GI denote gross income and A above-the-line adjustments. Adjusted gross income is:

$$AGI = GI - A.$$

Let D_s be the standard deduction and D_i itemized deductions. The deduction used is:

$$D = \max(D_s, D_i).$$

Taxable income is:

$$TI = \max(0, AGI - D - Q),$$

where Q is the QBI deduction. Total tax can be represented as:

$$T = \max(T_o + T_k + T_{SE} + T_{NIIT}, T_{AMT}) - CR.$$

For a candidate strategy h, savings are:

$$S_h = T_0 - T_h.$$

The strategy is not selected solely by maximizing S_h. After computation, the system stress-tests the recommendation under adverse changes such as income variation, deduction limits, tax-rate changes, phase-out effects, liquidity constraints, and audit scrutiny. A simulated CPA review layer evaluates professional defensibility, documentation needs, audit risk, client suitability, and communication requirements. Finally, the TaxEvaluationAgent decides whether to accept the best strategy, continue searching, or stop after the maximum number of rounds.

The bounded vocabulary is essential because it defines what the model is allowed to propose. Instead of allowing free-form advice, the system restricts candidate strategies to recognized planning actions, permitted taxpayer attributes, defined income categories, and computable assumptions. This prevents the model from inventing deductions, misapplying credits, or mixing concepts that the engine cannot verify. The legal validity gate then performs a second filter by asking whether the strategy is even available under the synthetic rule set before the tax engine spends effort estimating its effect.

The review and evaluation stages make the architecture professional rather than merely computational. A strategy with strong simulated savings may still fail if it requires unrealistic documentation, creates a high audit-risk posture, conflicts with the client's liquidity needs, or depends on assumptions that would be difficult to explain. The CPA-style reviewer converts those concerns into structured objections. The TaxEvaluationAgent then weighs tax savings, legal validity, stress robustness, audit defensibility, and suitability together. If the best proposal fails, the reasons for failure are injected into the next generation round so that subsequent strategies avoid the same weaknesses.

4.3. Demonstrated Contribution

The main result is a governance-first prototype for AI-assisted tax planning. The system demonstrates that tax strategy discovery can be structured as an AI4S loop: propose a strategy, compute the outcome, test the risk, obtain professional review, and use failure context to guide the next round.

The key result is separation of roles. The LLM does not compute tax liability through free-form reasoning. It does not decide legality. It does not authorize implementation. The deterministic `compute_tax()` engine calculates the numerical outcome, the validity gate filters unsupported strategies, and the CPA-style reviewer challenges the recommendation before it can be treated as viable.

The prototype also shows why audit governance must be embedded early. If an idea is legally unavailable or documentation-fragile, it should not become attractive merely because it produces simulated savings. The system therefore prevents persuasive but unsupported AI output from being mistaken for tax advice.

The prototype's strongest result is that it changes the unit of analysis from an AI answer to a governed tax-planning record. Each candidate can be traced from prompt, through structured strategy representation, through validity screening, through tax calculation, through stress testing, through CPA-style review, and into a final evaluation decision. That traceability is central to tax work because the reasoning behind a position may matter as much as the numerical result. A recommendation that cannot be explained, reproduced, and documented is not professionally useful.

The system also demonstrates that failure can improve the search process. If a candidate fails because it is not legally available, the next prompt can narrow the strategy space. If it fails because the savings disappear under stress, the next cycle can request more robust planning. If it fails because documentation would be burdensome, the next round can favor simpler strategies. This transforms rejected outputs into structured learning. The result is not an automated tax adviser, but a prototype for disciplined tax discovery under professional supervision.

The result is also educational. By showing which strategies fail and why, the system helps users understand that tax planning is governed by constraints rather than by the largest simulated savings number. A rejected candidate can teach that eligibility, timing, substantiation, and audit posture are part of the economic result. This makes the prototype useful as a demonstration of accountable professional reasoning.

4.4. Limitations and Research Extensions

The most important improvement would be to replace the synthetic tax code with jurisdiction-specific rules, including statutory references, administrative guidance, case law, phase-outs, credits, entity-level issues, timing rules, anti-abuse doctrines, and state-tax interactions. A stronger taxpayer profile would include real-world complexity such as multiple entities, retirement accounts, carryforwards, charitable strategies, estate-planning considerations, AMT exposure, basis tracking, and liquidity constraints.

The CPA review layer could be strengthened with an audit-defense checklist, required-documentation standards, professional-signoff workflow, and risk scoring. The system should distinguish between strategies that are mathematically valid, administratively practical, and professionally defensible.

A production version would also require legal review, compliance controls, client consent, versioned assumptions, secure records, and clear non-reliance boundaries. The architecture could eventually support professional workflow, but only if the AI remains a hypothesis generator and not an authority for tax positions.

The computational layer would also need a richer treatment of time. Real tax planning is often intertemporal: accelerating income, deferring deductions, harvesting losses, making retirement contributions, changing entity elections, or timing charitable gifts can improve one year while worsening another. A production engine should therefore compare multi-year after-tax outcomes, liquidity effects, and risk of future reversal rather than reporting only a single-year liability reduction. It should also preserve the assumptions used for each projection so that a reviewer can understand exactly why the strategy works.

The governance layer should include formal engagement boundaries. Users should be told when the system is providing educational analysis, when professional review is required, and when a strategy cannot be implemented without separate advice. Each candidate should have a versioned record of inputs, rule sources, assumptions, reviewer notes, and final disposition. The LLM should never be able to override a legal gate, suppress an audit warning, or present an unreviewed strategy as advice. These controls would make the prototype more aligned with real professional practice.

A stronger version should also model uncertainty in taxpayer facts. Real clients may have incomplete records, changing income, uncertain business expenses, or unresolved entity questions. The system should flag where a strategy depends on facts that must be verified before analysis can continue. That would prevent the model from treating assumed facts as established facts and would improve audit defensibility.

4.5. Reference Materials

The companion paper and notebook for this domain are:

- **Paper:** AI4S TAX PLANNING.
- **Standard notebook:** COLAB TAX PLANNING.

5. AI4S Applied to Civil Litigation Strategy

5.1. Domain Problem

Litigation strategy is a discovery problem over facts, evidence, legal theories, precedent, ethics, cost, timing, settlement incentives, and client objectives. A legal theory can be creative and persuasive while still being unethical, weakly supported, economically irrational, procedurally unavailable, or unlikely to survive adverse precedent. The risk of AI in litigation is that it may generate confident legal narratives without sufficient grounding in evidence, procedural posture, jurisdiction, or professional responsibility.

At a doctoral legal-analysis standard, the decisive question is not whether a theory is rhetorically elegant. It is whether the theory can satisfy burdens of pleading, proof, admissibility, preservation, privilege, remedies, and professional responsibility at the relevant procedural stage. Litigation strategy is a dynamic decision problem under legal constraint. The system must therefore respect both doctrinal validity and procedural timing.

The problem is to design an AI system that can explore legal strategy without pretending to be a lawyer. The LLM may propose theories, but those theories must be screened ethically, tested against precedent, evaluated economically, stress-tested, and reviewed by an attorney before being treated as serious candidates. In this setting, governance is not a post-processing step. It is the condition that makes exploration permissible.

This is particularly important because litigation strategy is adversarial. A theory that looks strong in isolation may collapse when a missing document appears, when a hostile venue applies a stricter standard, when damages are difficult to prove, or when settlement economics dominate expected trial value. The architecture must therefore evaluate legal theories as fragile hypotheses rather than polished arguments.

The difficulty is that litigation strategy combines several forms of uncertainty at once. The factual record may be incomplete, evidence may be disputed, witnesses may be unreliable, damages may be hard to prove, and the procedural posture may change the value of a claim. At the same time, legal research imposes hierarchy: binding precedent matters more than persuasive authority, adverse authority cannot be ignored, and ethical duties constrain what arguments may be advanced. A language model that produces a polished theory without respecting these limits can mislead a user into believing that a weak position is stronger than it is.

The AI4S problem is therefore not to automate legal judgment, but to discipline legal hypothesis generation. Each theory must be expressed in structured form, tied to facts, checked against ethical rules, compared with precedent, valued economically, stressed under adverse assumptions, and reviewed by a professional. The architecture must also make rejected theories informative. If a theory fails because evidence is missing, precedent is hostile, or the expected litigation value is too low, that reason should guide the next search round.

5.2. Governed Architecture

The litigation architecture begins with a synthetic legal universe, including causes of action, ethical rules, jurisdictional assumptions, procedural constraints, and precedent records. A case profile defines the dispute, evidence inventory, damages categories, client goals, risk tolerance, timing constraints, and settlement posture. The LLM generates candidate legal theories in structured form, but each candidate must pass an ethical gate before it can be simulated.

The evidentiary gate is the central legal safeguard. It prevents the model from treating assertions as facts, facts as admissible evidence, or narrative coherence as legal sufficiency. A generated theory must identify the evidence on which it relies, the procedural vehicle through which the theory would be advanced, and the professional constraints that might limit the lawyer's ability to pursue it.

The precedent engine then evaluates analogous synthetic cases. Let $\mathcal{P}_h$ be the precedent set for legal theory h. If W_h is the number of plaintiff wins and S_h is the number of split verdicts, the baseline win probability is:

$$\pi_h^{(0)} = \frac{W_h + 0.5S_h}{|\mathcal{P}_h|}.$$

Evidence adjustments modify the probability:

$$\pi_h = \text{clip}\left(\pi_h^{(0)} + \Delta_{\text{support}} - \Delta_{\text{gap}}, 0.10, 0.95\right).$$

If $E[D_h]$ denotes expected damages, expected value is:

$$EV_h = \pi_h E[D_h].$$

Settlement economics are evaluated through:

$$NEV_h = EV_h - C_L,$$

where C_L is litigation cost.

The system also performs stress tests against adverse precedent, missing evidence, hostile venue assumptions, lower damages realization, and settlement pressure. A simulated senior attorney reviews the recommendation for ethical soundness, evidentiary support, litigation posture, client alignment, and settlement-versus-trial judgment. The LitigationEvaluationAgent then decides whether the process should stop, continue, or reject all current theories.

This sequence deliberately prevents numerical simulation from laundering an impermissible theory. A theory that fails the ethical gate does not proceed to expected-value analysis, because assigning a positive expected value to an unethical strategy would give it a false appearance of legitimacy. Similarly, a theory with inadequate evidentiary support must carry that weakness into the precedent and damages analysis. The system is therefore designed so that facts, ethics, and law constrain the computation rather than being decorative notes after the fact.

The evaluation agent gives the workflow memory. It records whether a theory failed because the legal basis was weak, the facts were underdeveloped, damages were speculative, settlement posture was unfavorable, or attorney review identified unacceptable risk. Those reasons are then used to condition the next prompt. For example, the next round may ask for a narrower cause of action, a theory requiring fewer disputed facts, a settlement-oriented approach, or an argument better aligned with available precedent. The result is an iterative strategy process that explores options while preserving professional boundaries.

5.3. Demonstrated Contribution

The main result is a disciplined attorney-in-the-loop legal strategy prototype. The system demonstrates that AI can broaden the search for legal theories while preserving ethical, evidentiary, procedural, and professional constraints. It does not allow an unethical theory to gain legitimacy through numerical analysis, and it does not allow expected value to override attorney judgment.

The result is important because it shows how legal reasoning can be made more auditable. Each theory is generated, screened, simulated, stressed, reviewed, and either accepted or rejected with reasons. Rather than treating AI output as a finished memorandum, the workflow treats it as a hypothesis that must survive a sequence of professional controls.

The prototype also demonstrates the value of failure-context learning. If a theory fails because evidence is missing, precedent is adverse, damages are speculative, or the ethical gate rejects it, that failure becomes structured memory for the next round. The system can then search for theories that avoid the same defects.

The prototype also shows how an AI-assisted litigation workflow can distinguish creativity from authorization. Creativity is valuable because it can surface alternative causes of action, defenses, settlement frames, or evidentiary themes that a team may want to examine. Authorization is different. A theory becomes actionable only after it survives ethical screening, precedent analysis, factual review, damages modeling, stress testing, and attorney judgment. This distinction is central to responsible use because litigation can impose costs on courts, clients, counterparties, and third parties.

Another result is that the system makes tradeoffs explicit. A theory may have a higher simulated win probability but lower damages, or a larger damages estimate but greater proof risk. A settlement path may produce lower upside but better risk-adjusted client value. By recording these dimensions separately, the architecture avoids collapsing legal strategy into a single number. The attorney reviewer can then challenge whether the recommendation reflects professional judgment, client objectives, and procedural reality. The system's output is therefore a structured decision record, not a substitute for counsel.

This is the central governance benefit. The workflow can support brainstorming without allowing brainstorming to become advice, filing strategy, or settlement instruction. It gives a legal team a structured way to compare theories, identify evidentiary gaps, and record professional objections. The AI contribution is useful precisely because it is bounded by gates that protect ethical duties and client-centered judgment.

5.4. Limitations and Research Extensions

A production version would need actual legal research integration, including statutes, procedural rules, binding precedent, negative-treatment checks, jurisdictional hierarchy, evidentiary admissibility, limitations periods, venue-specific standards, and current-law validation. The precedent engine would need richer factual-similarity scoring, procedural context, burden-of-proof modeling, and explicit treatment of adverse authority.

The damages model would also need to incorporate proof standards, collectability, litigation cost, insurance coverage, counterclaims, discovery burden, motion practice, appeal risk, and settlement dynamics. Expected value should never be reduced to a single number without explaining the assumptions and uncertainty around that number.

Most importantly, the attorney-review layer should be replaced by actual licensed attorney review. A real system would require privilege-preserving workflows, confidentiality protections, client-consent controls, audit logs, conflict checks, matter-specific authorization, and safeguards against unauthorized practice of law. The architecture can support legal professionals, but it must not represent itself as legal counsel.

The research layer would also need source-quality controls. A production system should distinguish binding authority from persuasive authority, identify negative treatment, account for dates and jurisdictional hierarchy, and preserve citations to the exact legal materials used. It should not rely on a model's memory of law. Current-law validation is essential because rules, standards, and procedural requirements change. The system should therefore connect every simulated theory to verifiable sources and flag uncertainty where legal authority is incomplete or contested.

The workflow should also support matter management. Real litigation strategy depends on privilege, confidentiality, deadlines, discovery obligations, client budgets, settlement authority, and conflicts. A mature system would need role-based access, secure evidence handling, redaction controls, preservation of attorney work product, and clear logs showing who reviewed each recommendation. It should also include safeguards that prevent non-lawyers from using outputs as legal advice without counsel. These improvements would preserve the prototype's core value while making it safer for supervised legal teams.

The system should further distinguish procedural stages. A theory that is plausible at intake may not be viable after discovery, summary judgment, trial preparation, or appeal. Production use would therefore require stage-specific standards for evidence, cost, settlement leverage, and client authority. Treating procedural posture as a first-class input would make the simulation more faithful to actual litigation practice.

5.5. Reference Materials

The companion paper and notebook for this domain are:

- **Paper:** AI4S LITIGATION.
- **Standard notebook:** COLAB LITIGATION.

6. AI4S Applied to Financial Advice

6.1. Domain Problem

Financial advice is not simply portfolio optimization. A household has ranked goals, emotional constraints, liquidity needs, tax considerations, risk tolerance, behavioral tendencies, and changing life circumstances. A strategy that maximizes expected wealth may still be unsuitable if it jeopardizes retirement security, creates panic risk, ignores near-term liquidity needs, or misallocates resources across goals.

At a combined finance-and-law standard, suitability is not a soft preference layered on top of portfolio theory. It is the governing constraint that determines whether a recommendation can be made to a particular client. The household's financial life contains legal obligations, tax consequences, insurance needs, debt covenants, family commitments, and behavioral limits that cannot be reduced to expected return.

The problem is to design an AI system that supports financial planning without becoming a financial advisor. The LLM can propose planning strategies, but those strategies must be computed through simulation, screened for suitability, tested under stress scenarios, and reviewed by an advisor before being accepted. The system must respect the distinction between a mathematically attractive allocation and a client-appropriate recommendation.

This distinction is central because households do not experience risk as an abstract variance term. They experience risk as delayed retirement, missed college funding, forced asset sales, anxiety during drawdowns, or inability to absorb emergencies. The architecture must therefore evaluate strategies through the client's goal hierarchy and behavioral constraints, not merely through expected return.

The planning problem is therefore multi-objective and client-specific. Retirement security may deserve priority over discretionary spending, college funding may compete with debt reduction, and a tax-efficient portfolio may still be inappropriate if it creates complexity the household cannot manage. Liquidity is especially important because clients experience risk through bills, emergencies, job loss, health shocks, and behavioral pressure, not merely through portfolio variance. An AI-generated recommendation that ignores those lived constraints can be mathematically elegant and professionally unsuitable at the same time.

The AI4S framing converts a proposed financial plan into a testable planning hypothesis. The hypothesis must specify allocation, savings priorities, account usage, withdrawal or contribution behavior, risk assumptions, and the goals it is intended to support. It must then be computed by simulation, screened for suitability, tested under adverse scenarios, and reviewed through an advisor-style lens. The objective is not to let the model choose a client's financial life. The objective is to help surface candidate plans that can be examined, challenged, and either improved or rejected by a governed process.

This makes suitability the core problem rather than an afterthought. The system must recognize that clients do not experience financial plans as equations; they experience them through uncertainty, tradeoffs, habits, family obligations, and fear of irreversible mistakes. A responsible AI architecture must therefore test whether the proposal fits the client, not merely whether the simulated portfolio performs well.

6.2. Governed Architecture

The financial-advice architecture begins with a synthetic market universe and a client profile. The client profile includes income, portfolio holdings, savings capacity, risk tolerance, liquidity needs, known portfolio issues, and ranked goals such as retirement, college funding, and a vacation-home down payment. The LLM proposes candidate wealth strategies, including allocations, product mixes, contribution priorities, and account recommendations.

The suitability gate should therefore ask whether the plan is coherent for the client, not merely efficient in the model. A plan that improves expected terminal wealth may still fail if it concentrates risk, reduces emergency reserves, creates tax complexity, ignores insurance exposure, or requires behavior the household is unlikely to sustain. This is why the architecture treats client comprehension and implementation realism as part of the recommendation itself.

A suitability gate filters inappropriate strategies before simulation. A candidate strategy h defines allocation weights $\mathbf{w}^{(h)}$. Portfolio expected return is:

$$\mu_p^{(h)} = \sum_i w_i^{(h)} \mu_i.$$

Portfolio variance is:

$$\sigma_p^{2(h)} = \sum_i (w_i^{(h)})^2 \sigma_i^2 + 2 \sum_{i<j} w_i^{(h)} w_j^{(h)} \rho_{ij} \sigma_i \sigma_j.$$

For each goal g, with target T_g and horizon H_g, the inflated target is:

$$\tilde{T}_g = T_g (1 + \iota)^{H_g}.$$

The probability of reaching the goal is:

$$P_g^{(h)} = \frac{1}{N} \sum_{s=1}^{N} \mathbf{1} \left[V_{s,H_g}^{(h)} \geq \tilde{T}_g \right].$$

The goal-weighted score is:

$$G^{(h)} = \sum_g \omega_g P_g^{(h)}.$$

The system then stress-tests each plan under a market crash, inflation spike, sequence-of-returns shock, and liquidity pressure. A simulated CFP-style advisor reviews the strategy for fiduciary suitability, behavioral realism, fee justification, goal hierarchy, and client communication. The WealthEvaluationAgent decides whether the best strategy is strong enough to stop or whether the system should continue searching.

The goal hierarchy is the central design choice. Rather than treating terminal wealth as the only objective, the system evaluates whether the plan supports each stated goal and then weights those goals according to client priority. This matters because a plan that maximizes expected wealth by taking high equity risk may reduce the probability of meeting a near-term down-payment goal or may create drawdowns that the client cannot tolerate. The suitability gate therefore checks not only asset allocation, but also liquidity, concentration, complexity, product fit, and consistency with the stated risk profile.

The advisor-review layer adds qualitative discipline to the quantitative engine. Monte Carlo output can show distributions, shortfall probabilities, and stress outcomes, but it cannot by itself determine whether a household understands the plan, can maintain savings behavior, or will abandon the strategy in a downturn. The CFP-style reviewer evaluates those issues and flags where the recommendation needs explanation, simplification, or rejection. The WealthEvaluationAgent then decides whether the plan is good enough to stop the search or whether failure context should guide another generation round.

6.3. Demonstrated Contribution

The main result is a goal-aware, advisor-governed financial-planning architecture. The system demonstrates that AI-generated plans must be evaluated against client-specific objectives rather than generic return metrics. Retirement security, college funding, emergency liquidity, and discretionary goals are not interchangeable; the architecture reflects that hierarchy by scoring each goal separately and then combining them through explicit weights.

The system also shows that behavioral review matters. A plan that looks mathematically attractive may fail if the client cannot tolerate drawdowns, understand the recommendation, maintain contributions, or remain invested during stress. The advisor review layer therefore becomes a professional gate, not a decorative explanation layer.

The prototype further demonstrates how AI can broaden planning options without weakening fiduciary controls. The LLM may suggest alternatives, but suitability screening, Monte Carlo analysis, stress testing, and advisor review determine whether those alternatives are appropriate for the household.

The prototype's important contribution is that it evaluates advice in the language of goals rather than products. The model may propose allocations, contribution priorities, or account recommendations, but the system asks whether those choices improve the client's actual planning problem. This prevents the analysis from becoming a generic portfolio exercise. It also makes the tradeoffs visible: increasing retirement contributions may reduce near-term liquidity, taking more risk may improve expected wealth while worsening behavioral resilience, and prioritizing one goal may delay another.

The system also demonstrates that professional review is not a ceremonial step. The simulated advisor can reject a plan that appears strong in Monte Carlo terms if it is unsuitable for the client's temperament, time horizon, liquidity needs, or ability to implement. That review protects against a central risk of AI financial tools: the tendency to produce confident, personalized-sounding recommendations without a fiduciary process. By requiring suitability screening, stress testing, and advisor challenge, the architecture turns AI output into an input for supervised planning rather than advice itself.

The result is a planning record that can be discussed, corrected, and improved. Instead of receiving a single confident recommendation, the user receives a trace of assumptions, simulations, stress failures, suitability flags, and advisor concerns. That record supports better professional conversations because it shows where the recommendation is strong, where it is fragile, and what client information would be needed before implementation.

6.4. Limitations and Research Extensions

The main improvement would be to move from a simplified synthetic market universe to an account-aware, tax-aware, household-level planning engine. A stronger version would include Social Security, pensions, insurance, liabilities, tax lots, Roth conversion analysis, estate planning, college inflation, healthcare shocks, retirement spending policies, required minimum distributions, and estate-transfer goals.

The Monte Carlo engine should include rebalancing rules, withdrawal strategies, inflation regimes, fat tails, historical stress periods, dynamic contribution behavior, spending flexibility, and changing risk capacity over time. It should also distinguish between probability of goal success, magnitude of shortfall, liquidity resilience, tax efficiency, and behavioral sustainability.

The advisor-review layer should be connected to real client discovery, investment policy statements, compliance review, client communication records, and documented consent. A production version would need supervision, records retention, fee disclosure, product-governance controls, and clear boundaries that prevent AI output from being treated as personalized advice without professional approval.

A production planning engine should also become explicitly account-aware. The same asset allocation can have very different consequences depending on whether assets sit in taxable accounts, retirement accounts, education accounts, or employer plans. Tax lots, contribution limits, required distributions, withdrawal sequencing, insurance coverage, debt terms, and beneficiary designations can materially change the suitability of a recommendation. Adding these features would make the simulation closer to household planning instead of abstract portfolio modeling.

The governance layer should include documentation that a real adviser would need: client discovery notes, risk-tolerance evidence, investment policy statements, fee disclosures, rationale for recommendations, rejected alternatives, and client communication records. The system should also distinguish education from advice and require professional signoff before any recommendation is implemented. Over time, the evaluation agent could learn from plan reviews, client behavior, and goal updates, but it should remain bounded by compliance controls and advisor judgment rather than becoming an autonomous recommendation engine.

The system should also incorporate ongoing monitoring. A plan that is suitable today may become unsuitable after a job change, market crash, inheritance, health event, tax-law change, or shift in client goals. A mature version would periodically refresh assumptions, rerun stress tests, and route material changes back through advisor review. This would turn the architecture from a one-time planner into a governed planning lifecycle.

6.5. Reference Materials

The companion paper and notebook for this domain are:

- **Paper:** AI4S FINANCIAL ADVICE.
- **Standard notebook:** COLAB FINANCIAL ADVICE.

7. AI4S Applied to Investment Banking and M&A

7.1. Domain Problem

M&A target discovery is a complex professional search problem. The right acquisition target must fit strategy, valuation, financing capacity, regulatory feasibility, synergy defensibility, EPS accretion, integration complexity, and board governance. A target can look attractive strategically but fail because the price is too high, synergies are unrealistic, leverage exceeds comfort, regulatory risk is unacceptable, or the board story is weak.

At a senior finance and corporate-law standard, the relevant object is not a persuasive deal story. It is a board-reviewable transaction record. That record must connect strategic rationale to valuation evidence, financing capacity, diligence findings, process integrity, conflicts analysis, fiduciary context, and downside scenarios. The AI system must therefore support disciplined target screening without creating premature deal momentum.

The problem is to use AI for target discovery without allowing the model to become an M&A committee. The LLM can propose target shortlists and strategic theses, but valuation, synergy modeling, financing feasibility, accretion, stress testing, and board readiness must be governed by deterministic engines and senior banker review. In this setting, the model's useful role is to broaden search and articulate hypotheses; it should not authorize a transaction narrative.

This separation is essential because M&A decisions combine analysis, negotiation, fiduciary process, diligence, financing markets, and institutional judgment. A persuasive acquisition thesis can create momentum before the numbers, risks, and governance record are ready. The architecture therefore treats each target as a hypothesis that must earn its way into the board discussion.

The search problem is also path-dependent. A target that appears attractive in an initial screen may become unattractive after diligence, market reaction, financing feedback, regulatory analysis, or integration planning. Conversely, a target with a modest headline score may be strategically valuable if it fills a capability gap, accelerates product distribution, or blocks a competitor. The system must therefore avoid treating target discovery as a simple ranking exercise. It needs to expose the assumptions behind each thesis and make clear which risks would have to be resolved before a board could rely on the recommendation.

The AI4S framing makes each target a transaction hypothesis. The model may suggest why the asset could matter, but the hypothesis must be tested through valuation, synergy economics, financing structure, accretion, stress cases, and senior review. The purpose is not to replace bankers, corporate development teams, counsel, or boards. The purpose is to make early-stage target discovery more disciplined, auditable, and skeptical before organizational momentum forms around an untested deal story.

This is why the first screen must remain skeptical even when a target seems strategically elegant. The architecture should force the acquirer to name the value-creation mechanism, the diligence questions that could disprove it, and the governance approvals needed before momentum builds around the transaction.

7.2. Governed Architecture

The M&A architecture begins with a synthetic deal universe, an acquirer profile, board constraints, financing assumptions, and a target universe. The LLM generates a shortlist of acquisition candidates from the approved target universe. A regulatory screen evaluates antitrust, concentration, and approval-risk concerns before the target advances into financial modeling.

The deal-review gate is designed to separate screening utility from transaction authority. A target can advance only if the thesis is supported by valuation, financing, synergy, and governance evidence. The system must also preserve kill criteria: regulatory blocks, unaffordable leverage, unsupported synergies, diligence red flags, conflicts, or board-process defects should stop the workflow unless senior professionals record a reason to continue.

Deterministic engines then compute valuation ranges, synergy NPV, financing mix, pro forma leverage, EPS accretion or dilution, and deal score. For target i, valuation is based on EBITDA and revenue multiples:

$$EV_i^{EBITDA} = m_s^{EBITDA} E_i,$$

$$EV_i^{REV} = m_s^{REV} R_i.$$

A blended valuation estimate is:

$$EV_i^{mid} = \alpha EV_i^{EBITDA} + (1-\alpha) EV_i^{REV}.$$

Synergy value is modeled as:

$$S_i^{annual} = C_i + Y_i + F_i,$$

where C_i is cost synergy, Y_i is revenue synergy, and F_i is financial synergy. The NPV of synergies is:

$$NPV(S_i) = \sum_{t=1}^{H} \frac{S_{i,t}^{annual} q_i}{(1+r)^t}.$$

The financing condition is:

$$P_i = Cash_i + Debt_i + Stock_i.$$

Pro forma leverage is:

$$\frac{ND_i^{pro}}{EBITDA_i^{pro}} \leq 3.5.$$

EPS accretion is:

$$Accretion_i = \frac{EPS_i^{pro} - EPS_{acq}}{EPS_{acq}}.$$

The system selects the strongest candidate by score, subject to strategic, regulatory, leverage, and accretion constraints. Stress tests then examine integration failure, leverage pressure, higher financing costs, lower synergy capture, and competing-bid scenarios. A simulated Managing Director reviews the recommendation for board readiness, synergy defensibility, negotiation posture, diligence red flags, and presentation quality. The DealEvaluationAgent decides whether the system should stop, continue, or reject the current target set.

The scoring process is intentionally constrained. A target cannot advance merely because the LLM describes a persuasive strategic fit. It must pass through an approved universe, satisfy the regulatory screen, support a plausible purchase price, fit within financing capacity, and produce an acceptable risk-adjusted deal score. The deterministic engines make the economics explicit by separating enterprise value, synergy value, financing mix, leverage, and accretion. This separation allows reviewers to see whether the recommendation depends mainly on strategic fit, aggressive synergies, favorable financing, or optimistic valuation assumptions.

The stress and MD-review stages are especially important because M&A base cases often embed optimism. The system therefore asks what happens if integration is delayed, synergy capture is lower, financing costs rise, a competitor bids, or leverage remains elevated for longer than expected. The simulated Managing Director then evaluates whether the transaction is explainable to a board, whether the negotiation posture is credible, and whether diligence red flags have been surfaced. The DealEvaluationAgent uses this record to continue search, reject the target, or stop with a provisional recommendation.

7.3. Demonstrated Contribution

The main result is a board-governed M&A discovery prototype. It shows that AI can support target ideation while leaving valuation and transaction judgment under controlled processes. The LLM is useful for expanding the search space and articulating strategic theses, but deterministic deal engines compute whether a transaction is financially plausible.

The architecture emphasizes synergy skepticism. Many M&A failures arise from unrealistic synergy assumptions, weak integration planning, or overconfident base cases. By decomposing cost, revenue, and financial synergies, stress-testing integration failure, and requiring MD review, the system prevents optimism from becoming a board recommendation without challenge.

The prototype also makes the transaction process more auditable. Each target can be traced from shortlist generation through regulatory screen, valuation, synergy model, financing structure, accretion analysis, stress result, MD review, and final evaluation decision. This creates a governed bridge between target ideation and board-ready recommendation.

A second result is that the prototype decomposes deal attractiveness into reviewable components. Instead of presenting one polished acquisition narrative, it forces the process to show the target universe, strategic rationale, regulatory screen, valuation range, synergy assumptions, financing structure, leverage impact, accretion outcome, stress result, and banker review. This is valuable because many deal failures are not caused by lack of ideas; they are caused by insufficient challenge to the assumptions that make the idea appear compelling.

The architecture also reinforces the difference between a target idea and a board recommendation. An LLM can help generate candidate names, compare strategic themes, and draft a preliminary thesis, but it cannot determine fairness, fiduciary process, diligence sufficiency, or negotiation strategy. Those judgments require human professionals and institutional controls. By making the evaluation agent depend on deterministic outputs and MD-style objections, the notebook demonstrates how AI can support early M&A discovery without collapsing the governance process that protects boards and shareholders.

The result is especially useful for early screening, where many possible targets compete for limited senior attention. The system can document why a target deserves additional diligence or why it should be dropped before consuming banker, management, and board bandwidth. That discipline helps reduce narrative bias, protects against premature deal enthusiasm, and keeps the transaction process tied to evidence.

7.4. Limitations and Research Extensions

A production version would need live market data, actual financial statements, quality-of-earnings analysis, trading comparables, precedent transactions, discounted cash flow, purchase accounting, tax structuring, financing commitments, regulatory counsel, diligence workflows, and board materials. It would also need secure handling of confidential information and controls around material nonpublic information.

The synergy model should include one-time integration costs, customer churn, employee retention, revenue dis-synergies, product-roadmap risk, systems migration, management distraction, and accountable synergy ownership. The financing model should capture market windows, rating-agency constraints, covenant capacity, bridge financing, refinancing risk, and shareholder reaction.

The MD review should be replaced by real banker review with committee controls, fairness-process documentation, conflict checks, legal review, and board-governance records. A mature version should also include post-close learning, allowing realized synergy capture, integration outcomes, and transaction performance to improve future deal screening.

The valuation architecture should also be expanded beyond simplified multiples. A stronger system would triangulate trading comparables, precedent transactions, discounted cash flow, sum-of-the-parts analysis, leveraged buyout feasibility, and scenario-weighted valuation. It should also track market windows, financing spreads, equity-market volatility, rating-agency considerations, covenant capacity, and shareholder reaction. Those factors can change whether a transaction that is strategically attractive is actually executable.

The process layer would need institutional controls around confidentiality and conflicts. Real M&A work involves material nonpublic information, wall-crossing procedures, restricted lists, fairness processes, board minutes, diligence workstreams, regulatory counsel, and negotiation strategy. A mature platform should preserve a secure record of assumptions, approvals, diligence findings, banker comments, legal issues, and board materials. Post-close learning should then compare promised synergies with realized integration outcomes so that future target screening becomes less dependent on optimistic base cases.

A stronger system should also model process risk. Deal success depends not only on valuation and synergies, but also on timing, counterpart willingness, confidentiality, employee retention, customer reaction, regulatory review, and board alignment. These factors should become explicit fields in the evaluation record. The platform should also support kill criteria, so that certain diligence findings automatically stop the process unless senior decision-makers document why continued work is justified.

The final improvement would be stronger learning after close. The system should compare underwriting assumptions with realized revenue, margin, retention, integration, and balance-sheet outcomes, then feed those lessons into future screening. That feedback loop would make the platform more consistent with the AI4S principle that discovery improves when failures and surprises become structured evidence.

7.5. Reference Materials

The companion paper and notebook for this domain are:

- **Paper:** AI4S INVESTMENT BANKING.
- **Standard notebook:** COLAB INVESTMENT BANKING.

8. Cross-Domain Synthesis

8.1. The Same Architecture, Five Different Risk Objects

The five domains differ in their professional risk object:

Domain	Risk Object	Professional Gate
Algorithmic Trading	False alpha and deployment risk	Research governance review
Tax Planning	Invalid or audit-fragile tax strategy	CPA-style review
Civil Litigation	Unethical or weak legal theory	Attorney review
Financial Advice	Unsuitable or behaviorally fragile plan	CFP-style review
Investment Banking / M&A	Overpriced, fragile, or non-board-ready deal	Managing Director review

The computational engines also differ, but the governance logic is the same. A model-generated candidate has no authority until it has been computed, stressed, reviewed, and approved by the evaluation agent.

The table also shows why the framework cannot be reduced to a single generic AI-control checklist. The professional gate is domain-specific because the relevant duty is domain-specific. Research governance asks whether an apparent trading edge is statistically and economically credible. CPA-style review asks whether a tax position is legally available and documentable. Attorney review asks whether a theory is ethical, procedurally viable, and supported by evidence. CFP-style review asks whether a plan is suitable for a household. Managing Director review asks whether a deal thesis is financeable, diligence-ready, and board-defensible. A high-quality AI4S platform must therefore combine a common architecture with specialized professional standards.

8.2. The Central Governance Rule

The central rule of the collection is:

AI may expand the search space, but it may not collapse the review process.

This rule protects against a common failure mode: using AI to accelerate output while weakening professional controls. In this collection, acceleration is acceptable only if traceability, computation, stress testing, and review are preserved.

The rule also expresses a legal and fiduciary intuition. Delegation of tasks is not the same as delegation of responsibility. A firm may use machines to search, summarize, simulate, and draft, but it cannot responsibly outsource the duties of supervision, candor, loyalty, competence, diligence, or board process to a generative model. The review process is therefore not administrative overhead. It is the mechanism by which professional accountability survives technological acceleration.

8.3. Failure as an Asset

In all five systems, failure is not discarded. It becomes structured memory. A rejected strategy, fragile stress result, weak legal theory, unsuitable plan, or non-board-ready deal is converted into failure context and injected into the next cycle. This is the most direct AI4S feature of the architecture.

$$\text{Failure} \rightarrow \text{Diagnosis} \rightarrow \text{Revised Search} \rightarrow \text{New Candidate}.$$

8.4. Why This Matters

The report demonstrates that AI governance should not be appended after output generation. Governance must be embedded into the workflow itself. In high-accountability domains, the architecture is the control system.

This matters because the institutional failure mode is often procedural rather than purely technical. A system may use a strong model and still be unsafe if it gives users an easy path from fluent output to action. Conversely, a less capable model may be safer and more useful if its outputs are bounded, checked, documented, and escalated appropriately. The quality of the architecture is therefore measured not only by the intelligence of the model, but by the discipline with which the institution absorbs or rejects the model's proposals.

9. Conclusion

9.1. What the Report Establishes

This report has consolidated five applications of AI for Science to financial and professional advisory domains. Across algorithmic trading, tax planning, litigation strategy, household financial advice, and investment banking / M&A, the same methodological claim appears: generative AI is most useful when it is embedded inside a governed discovery loop. The report does not argue that language models should become autonomous traders, tax planners, lawyers, financial advisers, or investment bankers. It argues that the creative and organizational power of these models can be valuable when it is constrained by computation, stress testing, professional review, and explicit governance.

The first conclusion is conceptual. AI4S should be understood less as a particular technology and more as a disciplined pattern of inquiry. In the natural sciences, this pattern appears when models propose molecules, structures, mechanisms, or experiments that are then tested against evidence. In finance and advisory work, the same pattern appears when models propose strategies that must be translated into structured hypotheses, evaluated by deterministic engines, challenged by scenario analysis, and reviewed against professional standards. The object of discovery changes, but the logic remains recognizable. A proposal becomes meaningful only when it can be represented, tested, criticized, and revised.

The second conclusion is architectural. A high-accountability AI workflow cannot rely on the model's fluency as evidence of correctness. The architecture must separate the functions of generation, computation, stress testing, review, and decision. Each function limits a different failure mode. Generation expands the search space. Computation prevents vague narrative from substituting for measurable consequence. Stress testing reveals dependence on fragile assumptions. Professional review introduces domain skepticism. Evaluation governs whether the system should continue, stop, or halt. Together, these functions convert a conversational model into one component of a larger discovery system.

The third conclusion concerns accountability. In professional settings, responsibility cannot be assigned to an opaque model output. Clients, firms, courts, regulators, boards, and markets require reasons, records, standards, and reviewable decisions. The AI4S architecture responds by making the process auditable. A serious system should preserve prompts, structured candidates, parser decisions, assumptions, computations, stress results, reviewer objections, approvals, rejections, and failure context. This record does not eliminate risk, but it changes the nature of the risk. It makes the reasoning process inspectable rather than hidden inside a persuasive paragraph.

9.2. Lessons Across the Five Domains

The algorithmic-trading application shows why AI-generated ideas must become bounded research hypotheses before they matter. A trading story is not an alpha signal. A signal must specify data, timing, construction, position rules, transaction costs, and rejection criteria. Once translated into code, it can be backtested, stressed, and discarded if it fails. This is a useful discipline because most trading ideas should fail. The value of the system is not that it magically finds profitable strategies on demand. The value is that it can search broadly while preserving the research controls that prevent overfitting, narrative bias, and premature deployment.

The tax-planning application shows that computation alone is not enough. A tax strategy may appear economically attractive but still be inappropriate because of legality, documentation, audit exposure, client suitability, or professional ethics. The AI4S pattern is useful because it places a legal and CPA-style review layer between the generated idea and any recommendation. It treats tax planning as a governed search through permissible structures, not as a race to produce the largest modeled savings. This distinction is essential. In tax, a fragile strategy can create costs that do not appear in a simplified savings calculation.

The litigation application shows the same lesson in a different form. Legal strategy is not only a matter of rhetorical force. It depends on facts, evidence, procedure, precedent, ethics, cost, settlement dynamics, and client objectives. A model may generate a plausible theory, but plausibility does not make the theory responsible. The architecture therefore requires the candidate to be checked against evidentiary support, procedural posture, settlement exposure, and attorney-style critique. This preserves a central professional boundary: AI can assist with exploration, but it cannot authorize a legal position merely because it can express one persuasively.

The financial-advice application emphasizes suitability and human behavior. A plan can be mathematically coherent and still fail a household. Clients have liquidity needs, fears, habits, tax constraints, family responsibilities, time horizons, and asymmetric reactions to loss. A retirement recommendation, insurance decision, debt plan, or investment allocation must therefore be evaluated in relation to the client's life, not only in relation to an expected return. The AI4S structure helps by forcing generated recommendations through cash-flow analysis, risk testing, behavioral review, and CFP-style suitability gates. It makes advice more explicit about why a recommendation fits a particular person.

The investment-banking and M&A application highlights strategic and institutional complexity. A deal can look attractive in a headline valuation and still fail because the financing is weak, synergies are overstated, integration is unrealistic, regulatory risk is underestimated, or board materials do not support the recommendation. The AI4S pattern converts a transaction idea into a reviewable deal object: strategic rationale, valuation, financing, accretion-dilution, synergy case, diligence issues, process risk, and board-readiness. This is important because M&A is especially vulnerable to optimistic narratives. A governed workflow forces optimism to encounter evidence.

Across these domains, the shared lesson is that the professional risk object changes, but the governance grammar does not. False alpha, invalid tax strategy, weak litigation theory, unsuitable financial plan, and non-board-ready acquisition are different failures. Yet each failure can be made visible through the same sequence of candidate generation, structured parsing, deterministic evaluation, stress testing, professional critique, decision, and memory. This is the strongest evidence that the framework is not merely a collection of unrelated notebooks. It is a reusable pattern for accountable AI-assisted discovery.

9.3. Implications for Responsible Deployment

The practical implication is that organizations should not begin with the question, "How can AI make professional output faster?" They should begin with a different question: "What review process must never be weakened, and how can AI operate inside it?" Speed is valuable only when it does not remove the controls that give professional work its legitimacy. In high-stakes finance and advisory settings, the architecture must make it easier to inspect assumptions, reproduce calculations, identify conflicts, escalate uncertainty, and stop unsafe recommendations.

A mature implementation would require stronger controls than the demonstrations in this report. It would need secure data handling, access controls, confidentiality protections, versioned assumptions, model-risk review, compliance supervision, and clear human approval points. It would need parsers that reject unsupported fields, engines that expose calculation limits, stress tests that reflect realistic adverse conditions, and reviewers that can explain why a candidate fails. It would also need policies defining when AI assistance is allowed, when it is prohibited, and when additional professional signoff is required.

The system would also need to treat failure as an institutional asset. Many organizations record final recommendations but lose the reasoning behind rejected alternatives. AI4S suggests the opposite discipline. Failed candidates should be preserved with their failure reasons. A trading idea that fails because of turnover, a tax strategy that fails because of audit risk, a litigation theory that fails because of weak evidence, a financial plan that fails because of liquidity strain, and a deal thesis that fails because of integration fragility all contain useful information. When captured properly, failure context prevents the system from repeating the same mistake and helps guide the next search.

Human professionals remain central in this architecture. The point is not to reduce them to passive approvers of machine output. The point is to give them a more structured environment in which to exercise judgment. A professional can compare alternatives, interrogate assumptions, inspect stress results, challenge the model's framing, and document the basis for acceptance or rejection. This may improve professional judgment because it makes hidden assumptions more visible. It may also improve governance because it creates a record of how judgment was exercised.

The limits are equally important. The notebooks described in this report use synthetic data, simplified assumptions, and simulated review agents. They do not establish production readiness. They do not provide investment advice, tax advice, legal advice, financial advice, transaction advice, or any other professional recommendation. Real deployment would require independent validation, licensed professional review, regulatory analysis, cybersecurity controls, data governance, model monitoring, and institutional authorization. The contribution is therefore methodological rather than operational. It shows how such systems can be designed, not that they are ready to replace existing professional processes.

9.4. Future Direction

The next stage of work should focus on moving from illustrative notebooks to controlled research platforms. That transition would require richer domain ontologies, stronger parsing, more realistic data environments, and deeper integration with existing professional systems. In trading, this means connecting hypothesis generation to institutional research databases, execution-cost models, borrow constraints, factor-risk systems, and paper-trading records. In tax, it means linking planning candidates to current law, jurisdictional rules, document requirements, entity structures, audit history, and professional-signoff workflows. In litigation, it means connecting theory generation to pleadings, discovery records, admissible evidence, docket constraints, settlement history, and attorney-supervision protocols. In financial advice, it means integrating household data, planning software, portfolio systems, insurance records, tax assumptions, and client-consent processes. In M&A, it means connecting target screening to diligence workstreams, valuation databases, financing markets, regulatory review, integration plans, and board materials.

Future systems should also develop better evaluation standards. Accuracy alone is not enough because many professional recommendations are not single-label predictions. The more relevant questions are whether the candidate is complete, feasible, lawful, suitable, robust, explainable, and appropriately documented. Evaluation should therefore combine quantitative metrics with process metrics. A trading strategy can be measured by performance and robustness, but also by whether its assumptions are reproducible. A tax plan can be measured by projected savings, but also by legal support and audit posture. A litigation strategy can be measured by expected value, but also by evidentiary sufficiency and ethical compliance. A financial plan can be measured by goal attainment, but also by liquidity resilience and client comprehension. A transaction thesis can be measured by value creation, but also by diligence quality and board-readiness.

Finally, future work should examine how professionals actually interact with these systems. A well-designed AI4S platform should not encourage blind acceptance of generated output. It should invite challenge. It should make uncertainty visible, surface contrary evidence, show the sensitivity of conclusions to assumptions, and require professionals to document why they override warnings or approve fragile candidates. The human interface is therefore part of the governance architecture. If the interface hides assumptions, buries risks, or makes approval easier than review, the system will fail even if the underlying model is powerful. Responsible AI4S requires not only better models, but better professional workflows around those models.

9.5. Final Claim

The final claim is that finance is a demanding but productive setting for AI4S because finance exposes the difference between an answer and a governed recommendation. An answer can be generated quickly. A governed recommendation must survive representation, calculation, stress, review, documentation, and accountability. This distinction matters across all five domains. The model can propose an alpha signal, but the research process must decide whether it survives. The model can propose a tax plan, but legal and CPA-style gates must decide whether it is defensible. The model can propose a litigation theory, but evidence and ethics must constrain it. The model can propose a financial plan, but suitability must govern it. The model can propose a deal, but diligence, valuation, financing, and board process must test it.

If finance uses AI only to produce faster narratives, it will amplify the very risks that make the domain difficult: overconfidence, opacity, conflicts, model error, selective evidence, and weak documentation. If finance uses AI4S as a governed discovery architecture, the technology can serve a different purpose. It can broaden search while narrowing authorization. It can make alternatives visible while making approval harder. It can accelerate analysis while preserving the right to reject. It can transform failure from embarrassment into memory. These are not minor process details. They are the conditions under which powerful AI systems can be used responsibly in high-accountability professional work.

The broader lesson is that the future of AI in finance and advisory work should not be framed as automation without oversight. It should be framed as disciplined augmentation. The model proposes. The system tests. The professional reviews. The governance layer decides. The process learns from failure and repeats only when continued search is justified. That is the central architecture of AI4S in financial domains, and it is the reason the framework matters beyond any single notebook, model, or use case. AI may expand the search space, but it may not collapse the review process. In that rule lies the practical promise of AI4S for finance.

A. Repository and Companion Materials

The complete collection is hosted at:

https://alexdibol.github.io/ai4s_financial_domain

The repository includes papers, Colab notebooks, and workflow materials for all five domains:

1. AI4S Algorithmic Trading.
2. AI4S Tax Planning and Audit Governance.
3. AI4S Civil Litigation Strategy.
4. AI4S Financial Advice.
5. AI4S Investment Banking and M&A.

B. Non-Reliance Notice

This report is educational and methodological. It does not constitute investment advice, trading advice, tax advice, legal advice, accounting advice, financial advice, fiduciary advice, transaction advice, a fairness opinion, or a recommendation to buy, sell, trade, litigate, settle, structure, merge, acquire, finance, or implement any real-world action.

All examples rely on synthetic data, simplified assumptions, and simulated professional-review agents. Any adaptation to real-world use requires independent verification, professional review, legal and compliance approval, model-risk review, and institutional authorization.

Bibliography

[1] H. Wang *et al.*, "Scientific discovery in the age of artificial intelligence," *Nature*, vol. 620, pp. 47–60, 2023.

[2] J. Jumper *et al.*, "Highly accurate protein structure prediction with AlphaFold," *Nature*, vol. 596, pp. 583–589, 2021.

[3] H. Markowitz, "Portfolio selection," *The Journal of Finance*, vol. 7, no. 1, pp. 77–91, 1952.

[4] W. F. Sharpe, "Capital asset prices: A theory of market equilibrium under conditions of risk," *The Journal of Finance*, vol. 19, no. 3, pp. 425–442, 1964.

[5] A. Damodaran, *Investment Valuation: Tools and Techniques for Determining the Value of Any Asset*. Wiley, 2012.

[6] J. Rosenbaum and J. Pearl, *Investment Banking: Valuation, Leveraged Buyouts, and Mergers and Acquisitions*. Wiley, 2013.

[7] National Institute of Standards and Technology, *Artificial Intelligence Risk Management Framework*. NIST AI 100-1, 2023.

[8] S. Russell, *Human Compatible: Artificial Intelligence and the Problem of Control*. Penguin, 2020.